

Lecture 8: Markov Decision Processes

Lecturer: Steven Wu

Scribe: Kimberly Truong, Megan Li

8.1 Overview

We will start sequential decision making and reinforcement learning:

1. Elements of Markov Decision Processes (MDP) with a focus on finite horizons
2. Value functions: Bellman Equations/optimality
Note: Bellman is sometimes synonymous with dynamic programming
3. Value/policy iteration

8.2 MDP Notation

Consider the following notation:

- State space S
- Action space A
- Reward function $R : S \times A \rightarrow \Delta([0, 1])$, or $r \sim R(s, a)$
- Transition operator $P : S \times A \rightarrow \Delta(S)$
- Initial state distribution $\mu_0 \in \Delta(S)$
with initial state $s_0 \sim \mu_0$

Fact 1 *MDPs have the Markovian (“memorylessness”) property where its future state(s) are independent of its history.*

8.3 Finite Horizon

Remark 2 *Discounted and Infinite Horizon problems also exist, but we will focus on problems with finite horizons in this class.*

Consider the following notation:

- Horizon H
- Trajectory $\tau = (s_0, a_0, r_0, s_1, \dots, s_{H-1}, a_{H-1}, r_{H-1})$
- Policy π
- Objective

$$J(\pi) = \mathbb{E} \left[\sum_{h=0}^{H-1} r_h | S_0, a_{0:H-1} \sim \pi \right] \quad (8.1)$$

There can be as many as $|A|^{|S|^H}$ deterministic policies.

- Value function(s)
 - State-value function

$$V_h^\pi := \mathbb{E} \left[\sum_{h'=h}^{H-1} r_{h'} | s_h = s, a_{h:H-1} \sim \pi \right] \quad (8.2)$$

- Action-value function

$$Q_h^\pi := \mathbb{E} \left[\sum_{h'=h}^{H-1} r_{h'} | s_h = s, a_h = a, a_{h+1:H-1} \sim \pi \right] \quad (8.3)$$

8.3.1 (Non-stationary) Markov Policy

In general, a policy $\pi : \mathcal{H} \rightarrow \Delta(A)$, where \mathcal{H} represents all partial history.

Definition 1 *A **non-stationary Markov Policy** compresses all partial history into the current state and time step. π_h denotes a policy defined at time step h .*

$$\pi : S \times [H] \rightarrow \Delta(A)$$

8.3.2 Bellman Equations

For any policy π , we can derive the following *Bellman equations*:

$$V_h^\pi = \mathbb{E}[r_h + V_{h+1}^\pi(s_{h+1}) | s_h = s, a_h \sim \pi] \quad (8.4)$$

$$Q_h^\pi(s, a) = \mathbb{E}[r_h + Q_{h+1}^\pi(s_{h+1}, a_{h+1}) | s_h = s, a_h = a, a_{h+1} \sim \pi] \quad (8.5)$$

Remark 3 Considering just the first time step $h = 1, \pi = \pi_0$ yields:

$$J(\pi) = \mathbb{E}_{s_0 \sim \mu_0}[V^\pi(s_0)]$$

Theorem 4 Define $V^* = (V_0^*, \dots, V_H^*)$ recursively as

$$V_H^*(s) = 0 \quad (8.6)$$

$$\forall(s, h) : V_h^*(s) = \max_a \mathbb{E}_{s' \sim P(s, a)} [r(s, a) + V_{h+1}^*(s')] \quad (8.7)$$

Then, $\sup_{\pi: \mathcal{H} \rightarrow \Delta(A)} J(\pi) = \mathbb{E}_{s_0 \sim \mu}[V_0^*(s_0)]$. Now define the policy $\pi^* := (\pi_0^*, \dots, \pi_{H-1}^*)$ as

$$\forall(s, h) : \pi_h^*(s) = \arg \max_a \{\mathbb{E}_{s' \sim P(s, a)} [r(s, a) + V_{h+1}^*(s')]\} \quad (8.8)$$

Because π^* achieves value V^* for all (s, h) , it achieves optimality.

The equation (8.7) is called the Bellman optimality equation (for V). The recursive procedure (going from H to 0) defined by equations (8.6) and (8.7) is called *value iteration*.

Similarly, one can also derive the Bellman optimality equation and value iteration in terms of action value functions:

$$\forall(s, a) \in S \times A \quad Q_H^*(s, a) = 0 \quad (8.9)$$

$$\forall(s, a, h) \in S \times A \times [H - 1] \quad Q_h^*(s, a) = \mathbb{E} \left[r(s, a) + \max_{a'} Q_{h+1}^*(s', a') \right] \quad (8.10)$$

The optimal policy can also be written as

$$\pi_h^*(s) = \arg \max_a Q_h^*(s, a)$$

8.3.3 Policy Iteration

Another algorithm for computing the optimal value function and policy is *policy iteration*. Start with $\pi^{(0)}$, and then for each increment of t , we compute $Q^{\pi^{(t-1)}}$ where $Q_H^{\pi^{(t-1)}} = 0$ and for all (s, a) pairs, $Q_h^{\pi^{(t-1)}}(s, a)$ is computed by (8.5). This is the *policy evaluation* step. It is followed by the *policy improvement* step, during which we greedily update the policy as

$$\pi_h^{(t)}(s) := \arg \max_a Q_h^{\pi^{(t-1)}}(s, a) \quad (8.11)$$

Remark 5 *This algorithm guarantees local improvement—that is, at every time step h and state s :*

$$\forall(s, h) : \mathbb{E}_{a \sim \pi^{(t+1)}(s)}[Q_h^{\pi^{(t)}}(s, a)] = \max_{a \in A} Q_h^{\pi^{(t)}}(s, a) \geq \mathbb{E}_{a \sim \pi^{(t)}}[Q_h^{\pi^{(t)}}(s, a)] \quad (8.12)$$

for all t .

Interestingly, by the seminal *Performance Difference Lemma*, we can ensure *global improvement*, provided that we can local improvement at every state s . In this context, the lemma can be stated as follows.

Lemma 6 (Performance Difference Lemma)

$$J(\pi^{(t+1)}) - J(\pi^{(t)}) = \mathbb{E}_{\tau \sim \pi^{(t+1)}} \left[\sum_{h=1}^{H-1} \mathbb{E}[Q^{\pi^{(t)}}(s_h, \pi_h^{(t+1)}(s_h))] - \mathbb{E}[Q^{\pi^{(t)}}(s_h, \pi_h^{(t)}(s_h))] \right] \quad (8.13)$$

The left-hand side of the equation is the “global improvement” in the total reward objective and the right-hand side is the sum over the expected local improvement over time steps.